



Samueli
Computer Science



CS M146 Discussion: Week 5

Overfitting and Regularization, Neural Nets

Junheng Hao
Friday, 02/05/2021



Roadmap

- Announcement
- Overfitting and Regularization
- Neural Nets

Announcements



- **5:00 pm PST, Feb 5 (Friday):** Weekly Quiz 5 released on Gradescope.
- **11:59 pm PST, Feb 7 (Sunday):** Weekly quiz 5 closed on Gradescope!
 - Start the quiz before **11:00 pm PST, Feb 7** to have the full 60-minute time
- **Problem set 2** released on CCLE, submission on Gradescope.
 - Please assign pages of your submission with corresponding problem set outline items on GradeScope.
 - You do not need to submit code, only the results required by the problem set
 - Due on **11:59pm PST, Feb 12 (Friday)**






About Quiz 5

- Quiz release date and time: **Feb 5, 2021 (Friday) 05:00 PM PST**
- Quiz due/close date and time: **Feb 7, 2021 (Sunday) 11:59 PM PST**
- You will have up to **60 minutes** to take this exam. → Start before **11:00 PM** Sunday
- You can find the exam entry named "Week 4 Quiz" on GradeScope.
- Topics: **Overfitting, Regularization, Neural Nets (without Backprop)**
- Question Types
 - True/false, multiple choices
 - Some questions may include several subquestions.
- Some light calculations are expected. Some scratch paper and one scientific calculator (physical or online) are recommended for preparation.



Logistic regression (without regularization) **cannot** converge on a linearly separable dataset.



Clarification: Logistic Regression Convergence



```
X, y = np.array([[-1],[1]]), np.array([0,1]) # train data
```

In sklearn, you have solver options as `newton-cg`, `lbfgs`, `liblinear`, `sag`, `saga`.

Then train a logistic regression model **without** penalty

```
clf = LogisticRegression(random_state=0, penalty='none', solver='sag', max_iter=1000).fit(X, y) # or other solver except 'liblinear'
```

You may notice different solvers may result in different w ranging from 5 to 10 (printed by `clf.coef_`). Sometimes you might have a convergence error as follows:

```
clf=None
# solver = 'newton-cg', 'lbfgs', 'liblinear', 'sag', 'saga'
clf = LogisticRegression(random_state=12, penalty='none', solver='sag', max_iter=1000, verbose=10).fit(X, y)
clf.coef_, clf.intercept_, clf.n_iter_

max_iter reached after 0 seconds

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
/Users/junhenghao/opt/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_sag.py:330: ConvergenceWarning: The
max_iter was reached which means the coef_ did not converge
"the coef_ did not converge", ConvergenceWarning)
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 0.0s remaining: 0.0s
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 0.0s finished

(array([[8.29936548]]), array([0.00109003]), array([1000], dtype=int32))
```

And again train a logistic regression model **with** L2 penalty

```
clf = LogisticRegression(random_state=0, penalty='l2', solver='lbfgs').fit(X, y) # L2 used
```

This time, you will find different solvers converges to $w = 0.675$.



Overfitting

Key Questions:





- How to identify overfitting?
- How to avoid overfitting?

Credit:


<https://hackernoon.com/memorizing-is-not-learning-6-tricks-to-prevent-overfitting-in-machine-learning-820b091dc42>

	Low Training Error	High Training Error
Low Testing Error	The model is learning!	Probably some error in your code. Or you've created a <i>psychic AI</i> .
High Testing Error	OVERFITTING	The model is not learning.

Overfitting: Polynomial Regression




Overfitting




Overfitting: Data solution

- Collecting more data
- Data augmentation



Overfitting: Model Solution

- Avoid overfitting by changing model hyperparameter selection, from the mechanism and inductive bias of the model.





Regularization

Linear Regression

- Model

$$\hat{y} = \mathbf{x}^T \boldsymbol{\beta} = x_1\beta_1 + x_2\beta_2 + \cdots + x_p\beta_p$$

- Original Objective

$$\min_{\boldsymbol{\beta}} J(\boldsymbol{\beta}) = \frac{1}{2} \sum_{i=1}^N (\mathbf{x}^T \boldsymbol{\beta} - y)^2$$

- L2-Regularized Objective

$$\min_{\boldsymbol{\beta}} J(\boldsymbol{\beta}) = \frac{1}{2} \sum_{i=1}^N (\mathbf{x}^T \boldsymbol{\beta} - y)^2 + \boxed{\frac{\lambda}{2} \|\boldsymbol{\beta}\|^2}$$

Logistic Regression

- Model

$$y = \sigma(X) = \frac{1}{1 + e^{-X^T \boldsymbol{\beta}}}$$

- Original Objective

$$J(\boldsymbol{\beta}) = -\frac{1}{n} \sum_i (y_i \mathbf{x}_i^T \boldsymbol{\beta} - \log(1 + \exp\{\mathbf{x}_i^T \boldsymbol{\beta}\}))$$

- L2-Regularized Objective

$$J(\boldsymbol{\beta}) = -\frac{1}{n} \sum_i (y_i \mathbf{x}_i^T \boldsymbol{\beta} - \log(1 + \exp\{\mathbf{x}_i^T \boldsymbol{\beta}\})) + \boxed{\lambda \sum_j \beta_j^2}$$



$$\min_{\beta} J(\beta) = \frac{1}{2} \sum_{i=1}^N (\mathbf{x}^T \beta - y)^2 + \frac{\lambda}{2} \|\beta\|^2$$

$$\frac{\partial J(\beta)}{\partial \beta} = \sum_{i=1}^N \mathbf{x}(\mathbf{x}^T \beta - y) + \frac{\partial \lambda \|\beta\|^2}{\partial \beta}$$

$$\frac{\partial J(\beta)}{\partial \beta} = \sum_{i=1}^N \mathbf{x}(\mathbf{x}^T \beta - y) + \lambda \beta$$

$$\nabla_{\beta} J(\beta) = 0$$

$$\nabla_{\beta} \left(\frac{1}{2} \|\mathbf{X}\beta - y\|^2 + \frac{\lambda}{2} \|\beta\|^2 \right) = 0$$

$$\nabla_{\beta} \left(\frac{1}{2} (\mathbf{X}\beta - y)^T (\mathbf{X}\beta - y) + \frac{\lambda}{2} \beta^T \beta \right) = 0$$

$$\mathbf{X}^T \mathbf{X} \beta - \mathbf{X}^T y + \lambda I \beta = 0$$

$$\beta = (\mathbf{X}^T \mathbf{X} + \lambda I)^{-1} \mathbf{X}^T y$$

About Norms (Vectors)

$$\|\mathbf{w}\|_1 = |w_1| + |w_2| + \dots + |w_N|$$




1-norm (also known as L1 norm)

$$\|\mathbf{w}\|_2 = \left(|w_1|^2 + |w_2|^2 + \dots + |w_N|^2 \right)^{\frac{1}{2}}$$


2-norm (also known as L2 norm or Euclidean norm)

$$\|\mathbf{w}\|_p = \left(|w_1|^p + |w_2|^p + \dots + |w_N|^p \right)^{\frac{1}{p}}$$


p-norm



Baseline



L1 Regularization



L2 Regularization



About Norms (Vectors)

$$\|\mathbf{w}\|_1 = |w_1| + |w_2| + \dots + |w_N|$$

1-norm (also known as L1 norm)

$$Loss = Error(y, \hat{y}) + \lambda \sum_{i=1}^N |w_i|$$

Loss function with L1 regularisation

$$\|\mathbf{w}\|_2 = \left(|w_1|^2 + |w_2|^2 + \dots + |w_N|^2 \right)^{\frac{1}{2}}$$

2-norm (also known as L2 norm or Euclidean norm)

$$Loss = Error(y, \hat{y}) + \lambda \sum_{i=1}^N w_i^2$$

Loss function with L2 regularisation

How does L1/L2 regularization change the gradient descent step?



About "Entrywise" Norms (Matrix)

- L(2,1) Norm and L(p,q) Norm

$$\|A\|_{2,1} = \sum_{j=1}^n \|a_j\|_2 = \sum_{j=1}^n \left(\sum_{i=1}^m |a_{ij}|^2 \right)^{\frac{1}{2}} \quad \Rightarrow \quad \|A\|_{p,q} = \left(\sum_{j=1}^n \left(\sum_{i=1}^m |a_{ij}|^p \right)^{\frac{q}{p}} \right)^{\frac{1}{q}}$$

- Frobenius norm (Hilbert–Schmidt norm)

$$\|A\|_F = \sqrt{\sum_{i=1}^m \sum_{j=1}^n |a_{ij}|^2}$$

- Max Norm

$$\|A\|_{\max} = \max_{ij} |a_{ij}|.$$



Effect on Regularization

Without L2 Regularization

```
clf = LogisticRegression(random_state=0, penalty='none', solver='lbfgs', max_iter=100).fit(X, y) # or other solver except 'liblinear'  
print(clf.coef_, clf.intercept_, clf.n_iter_)
```

```
[[9.91926856]] [0.] [13]
```

```
clf = LogisticRegression(random_state=0, penalty='none', solver='newton-cg', max_iter=1000).fit(X, y) # or other solver except 'liblinear'  
print(clf.coef_, clf.intercept_, clf.n_iter_)
```

```
[[10.20283614]] [0.] [9]
```

With L2 Regularization


```
clf = LogisticRegression(random_state=0, penalty='l2', solver='lbfgs').fit(X, y)  
print(clf.coef_, clf.intercept_, clf.n_iter_)
```


```
[[0.67483169]] [0.] [4]
```

```
clf = LogisticRegression(random_state=0, penalty='l2', solver='newton-cg').fit(X, y)  
print(clf.coef_, clf.intercept_, clf.n_iter_)
```

```
[[0.67482829]] [0.] [2]
```

Neural Networks: Neuron/Perceptron





NN Example: XOR

- Which NN architecture corresponds to which function?

Y	1	0	1
	0	0	0
X	0	1	



Y	1	1	1
	0	0	1
X	0	1	

Table 1: Truth table for AND

Table 2: Truth table for OR

Y	1	1	0
	0	0	1
X	0	1	

Table 3: Truth Table for XOR



<https://datascience.stackexchange.com/questions/11589/creating-neural-net-for-xor-function>

<http://ven.cs.stir.ac.uk/~kit/techreps/pdf/TR148.pdf>

<https://medium.com/@jayeshbahire/the-xor-problem-in-neural-networks-50006411840b>



NN Example: XOR

	1	0	1
Y	0	0	0
X		0	1

Table 1: Truth table for AND

	1	1	1
Y	0	0	1
X		0	1

Table 2: Truth table for OR



	1	1	0
Y	0	0	1
X		0	1

Table 3: Truth Table for XOR

<https://datascience.stackexchange.com/questions/11589/creating-neural-net-for-xor-function>


<http://ven.cs.stir.ac.uk/~kit/techreps/pdf/TR148.pdf>

<https://medium.com/@jayeshbahire/the-xor-problem-in-neural-networks-50006411840b>




NN Example: XOR

Linear classifiers
cannot solve this



$$\begin{array}{ll} \sigma(20*0 + 20*0 - 10) \approx 0 & \sigma(-20*0 - 20*0 + 30) \approx 1 \\ \sigma(20*1 + 20*1 - 10) \approx 1 & \sigma(-20*1 - 20*1 + 30) \approx 0 \\ \sigma(20*0 + 20*1 - 10) \approx 1 & \sigma(-20*0 - 20*1 + 30) \approx 1 \\ \sigma(20*1 + 20*0 - 10) \approx 1 & \sigma(-20*1 - 20*0 + 30) \approx 1 \end{array}$$



We can also use the
step function
instead of sigmoid



NN Example: XOR (Cont'd)

Example: XOR

Now let's consider using a two-layer neural network, with the following equation:

$$g(\mathbf{x}) = \mathbf{w}^T \max(0, \mathbf{W}^T \mathbf{x} + \mathbf{c}) + b$$

We haven't yet discussed how to optimize these parameters, but the point here is to show that by introducing a simple nonlinearity like $f(x) = \max(0, x)$, we can now solve the xor(\cdot) problem. Consider the solution:

$$\mathbf{W} = \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix}$$


$$\mathbf{c} = [0, -1]^T$$

$$\mathbf{w} = [1, -2]^T$$

2-Layer NN Example


Neural network architecture

An example 2-layer network is shown below.



Here, the three dimensional inputs ($\mathbf{x} \in \mathbb{R}^3$) are processed into a four dimensional intermediate representation ($\mathbf{h} \in \mathbb{R}^4$), which are then transformed into the two dimensional outputs ($\mathbf{z} \in \mathbb{R}^2$).

2-Layer NN Example





- Layer 1: $\mathbf{h}_1 = f(\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1)$
- Layer 2: $\mathbf{h}_2 = f(\mathbf{W}_2 \mathbf{h}_1 + \mathbf{b}_2)$
- \vdots
- Layer N: $\mathbf{z} = \mathbf{W}_N \mathbf{h}_{N-1} + \mathbf{b}_N$

Questions:



1. Neural network model (in equations)
2. Number of neurons?
3. Number of weight parameters / bias parameters / total learnable parameters?

Neural Networks: Demo

- Let's play with it:
<https://playground.tensorflow.org/>



Neural Networks: Backpropagation



Neural Networks: Backpropagation


- A simple example to understand the intuition
- $f(x,y) = x^2y + y + 2$
- Forward pass:
 - $x=3, y=4 \rightarrow f(3,4)=42$

- Backward pass:
 - Chain rule:


$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial n_i} \times \frac{\partial n_i}{\partial x}$$

Another better demo:

<http://colah.github.io/posts/2015-08-Backprop/>




Multiclass Classification




Key: **sharing the same hidden layers** with **different weights** at the end

Question: Pros and cons?




2-Layer NN Example

Demo in class : Back propagation for a 2-layer network



Backprop: Exercise (next week)



In this question, let's consider a simple two-layer neural network and manually do the forward and backward pass. For simplicity, we assume our input data is two dimension. Then the model architecture looks like the following. Notice that in the example we saw in class, the bias term b was not explicit listed in the architecture diagram. Here we include the term b explicitly for each layer in the diagram. Recall the formula for computing $x^{(l)}$ in the l -th layer from $x^{(l-1)}$ in the $(l-1)$ -th layer is $x^{(l)} = f^{(l)}(\mathbf{W}^{(l)}x^{(l-1)} + \mathbf{b}^{(l)})$. The activation function $f^{(l)}$ we choose is the sigmoid function for all layers, i.e. $f^{(l)}(z) = \frac{1}{1+\exp(-z)}$. The final loss function is $\frac{1}{2}$ of the mean squared error loss, i.e. $l(\mathbf{y}, \hat{\mathbf{y}}) = \frac{1}{2}\|\mathbf{y} - \hat{\mathbf{y}}\|^2$.



We initialize our weights as

$$\mathbf{W}^{(1)} = \begin{bmatrix} 0.15 & 0.2 \\ 0.25 & 0.3 \end{bmatrix}, \quad \mathbf{W}^{(2)} = [0.4, 0.45], \quad \mathbf{b}^{(1)} = [0.35, 0.35], \quad \mathbf{b}^{(2)} = 0.6$$

- When the input $\mathbf{x}^{(0)} = [0.05, 0.1]$, what will be the value of $\mathbf{x}^{(1)}$ in the hidden layer? (Show your work).
- Based on the value $\mathbf{x}^{(1)}$ you computed, what will be the value of $\mathbf{x}^{(2)}$ in the output layer? (Show your work).
- When the target value of this input is $y = 0.01$, based on the value $\mathbf{x}^{(2)}$ you computed, what will be the loss? (Show your work).


Why understanding backpropagation?

- “Why do we have to write the backward pass when frameworks in the real world, such as TensorFlow/PyTorch, compute them for you automatically?”
- Vanishing gradients on Sigmoids




Why understanding backpropagation?

- ReLUs



- tanh





Why understanding backpropagation?

- Examples of activation function: Sigmoid, ReLU, leaky ReLU, **tanh**, etc
- Properties we focus:
 - Differentiable
 - Range: Whether saturated or not? (
 - Whether zero-centered or not?
- Activation function family
 - Wiki: https://en.wikipedia.org/wiki/Activation_function



- Backpropagation (CS 231N at Stanford)
 - <https://cs231n.github.io/optimization-2/>
 - <https://www.youtube.com/watch?v=i94OvYb6noo>
- (Optional) Matrix-Level Operation:
 - <https://medium.com/@14prakash/back-propagation-is-very-simple-who-made-it-complicated-97b794c97e5c>

NN: Number of iterations to converge

- Architecture/Meta-parameters of the network, e.g. # layers, activation
- Quality of training data (input-output correlation, normalization, noise cleansing, class distribution/imbalance)
- Random initialization of the parameters/weights
- Optimization algorithm, e.g. SGD, Adam, etc.
- Learning rate
- Batch size
- (In practice) Implementation quality (Bug-free? Optimized?)

<https://medium.com/datathings/neural-networks-and-backpropagation-explained-in-a-simple-way-f540a3611f5e>

<https://www.quora.com/Machine-Learning-What-are-some-tips-and-tricks-for-training-deep-neural-networks>


NN Summary




A mostly complete chart of Neural Networks

©2016 Fjodor van Veen - asimovinstitute.org


Deep Feed Forward (DFF)




Single Layer (Perceptron)



1 Hidden Layer




2 Hidden Layers



Flexibility

- https://www.packtpub.com/mapt/book/big_data_and_business_intelligence/9781788397872/1/ch01lvl1sec27/pros-and-cons-of-neural-networks
- <http://kseow.com/nn/>
- <https://towardsdatascience.com/hype-disadvantages-of-neural-networks-6af04904ba5b>

NN Summary: Pros and Cons




Efficiency (In many cases, prediction/inference/testing is fast)

https://www.packtpub.com/mapt/book/big_data_and_business_intelligence/9781788397872/1/ch01lvl1sec27/pros-and-cons-of-neural-networks


<http://www.luigifreda.com/2017/04/08/cnn-slam-real-time-dense-monocular-slam-learned-depth-prediction/>

<http://www.missqt.com/google-translate-app-now-supports-instant-voice-and-visual-translations/>

NN Summary: Pros and Cons



We trained both our baseline models for about 600,000 iterations (33 epochs) – this is similar to the 35 epochs required by Nallapati et al.'s (2016) best model. Training took 4 days and 14 hours for the 50k vocabulary model, and 8 days 21 hours for the 150k vocabulary model. We found the pointer-generator model quicker to train, requiring less than 230,000 training iterations (12.8 epochs); a total of 3 days and 4 hours. In particular, the pointer-generator model makes much quicker progress in the early phases of training. This work was begun while the first author was an intern at Google Brain and continued at Stanford. Stanford University gratefully acknowl-





Efficiency (Big model → slow training, huge energy consumption (e.g. for cell phone))

<https://www.kdnuggets.com/2017/08/first-steps-learning-deep-learning-image-classification-keras.html/2>

See, Abigail, Peter J. Liu, and Christopher D. Manning. "Get to the point: Summarization with pointer-generator networks." *arXiv preprint arXiv:1704.04368* (2017).


<https://www.lifewire.com/my-iphone-wont-charge-what-do-i-do-2000147>

NN Summary: Pros and Cons



Data (Both a pro and a con)

NN Summary: Pros and Cons



Computational Power (Both a pro and a con)


<https://www.anandtech.com/show/10864/discrete-desktop-gpu-market-trends-q3-2016>

<https://www.zdnet.com/article/gpu-killer-google-reveals-just-how-powerful-its-tpu2-chip-really-is/>

NN Summary: Pros and Cons



Black Box
Interpretability



What's next?



- In next week's discussion, we will continue to discuss:
 - Backpropagation in neural nets
- Programming Guide
 - PyTorch (for PS3)



Samueli
Computer Science



Thank you!

Q & A